

Data Preparation: Cleaning and Wrangling

Bloque III

Data preparation: Cleaning and Wrangling

- The context of data analytics
 - One of the most common mistakes of data analytic projects is thinking that they start with “analysis”
 - In their natural form, the “Raw Data” usually have registry errors that make an exact analysis impossible
 - All the data records have to be pre-processed:
 - In other words, the data must be cleaned, unified, consolidated and normalized, so that it can be used and extract valuable information.




rs743534	10	135349226	GG	DNA raw data
rs7435347	4	156654735	AG	
rs743535	10	135349367	CC	
rs743537	9	130634944	AA	
rs743539	9	130633904	TG	
rs7435399	4	24765819	TC	
rs7435410	4	89456219	AG	
rs7435411	4	169426096	GA	
rs7435420	4	62150059	GA	
rs743554	17	73754248	CC	
rs7435595	4	179402198	AG	
rs7435615	4	20399168	AG	
rs743562	5	131872383	CT	
rs7435646	4	100237492	TA	
rs743567	14	23890982	AA	
rs7435692	4	88489929	AC	
rs743572	10	104597152	AG	
rs7435768	4	116426581	TC	
rs743577	1	11119135	GA	

Genomapp

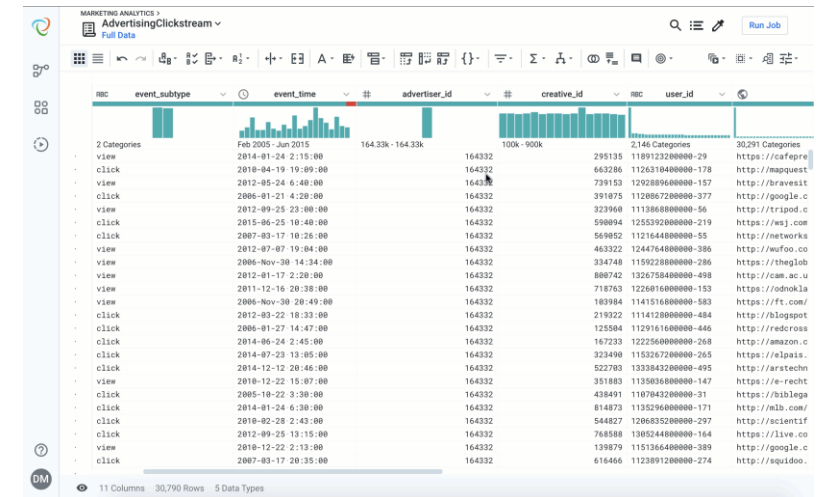
Data preparation: Cleaning and Wrangling

- The context of data analytics
 - Any data project needs a previous step to be successful, which comprises two main tasks:
 - Data Cleaning
 - Data Wrangling
 - In fact, they are usually the most time consuming for data analysts
 - According to a survey conducted in 2017, a data analyst can spend, on average, **80% of their time** on Data Wrangling.

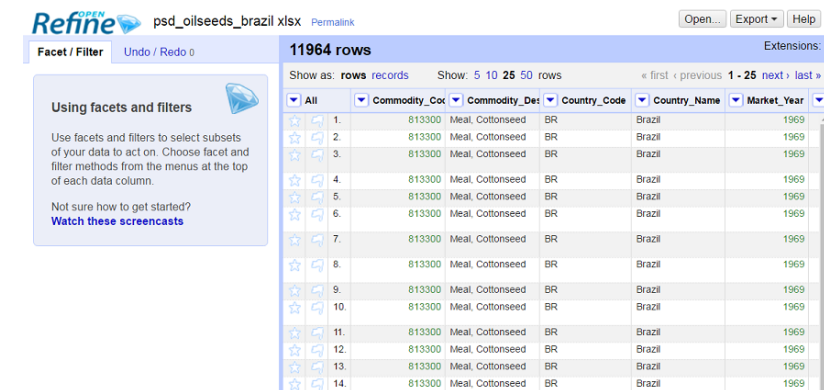


Data preparation: Cleaning and Wrangling

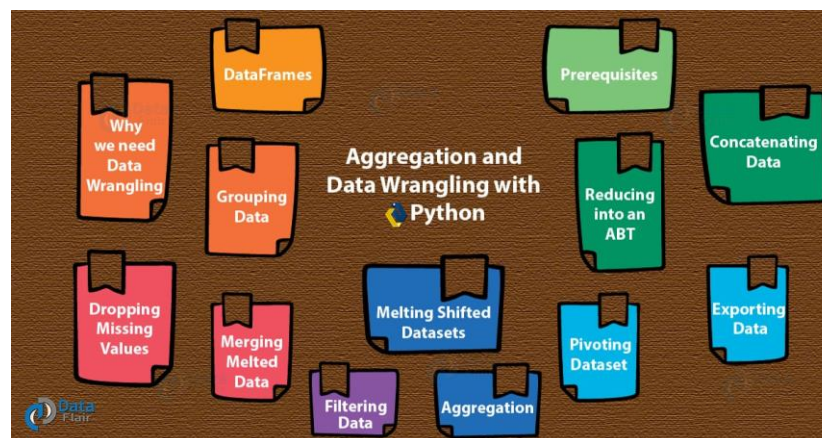
- Tools and methods for data preparation
 - Commercial: Trifacta <https://www.trifacta.com/>
 - Open source: Open Refine <https://openrefine.org/>
- Do it by yourself!



event_subtype	event_time	advertiser_id	creative_id	user_id
view	2014-01-24 2:15:00	164332	100k-900k	2146 Categories
click	2018-04-19 19:09:00	164332	295135	1189123200000-29
view	2012-05-24 6:40:00	164332	663286	1126310400000-178
click	2006-01-21 4:20:00	164332	739153	1292889600000-157
view	2012-09-25 23:00:00	164332	391875	1128672000000-377
click	2015-06-25 10:40:00	164332	323968	1113868800000-56
view	2007-03-17 10:26:00	164332	598094	1255392000000-219
click	2012-07-07 19:04:00	164332	569852	1121640000000-55
view	2006-Nov-30 14:34:00	164332	463322	1247640000000-386
click	2012-01-17 2:20:00	164332	334748	1159228800000-286
view	2011-12-16 20:38:00	164332	880742	1326758400000-498
click	2006-Nov-30 20:49:00	164332	718763	1226816000000-153
view	2012-09-22 18:33:00	164332	183984	1141516800000-583
click	2006-01-27 14:47:00	164332	219322	1141200000000-484
view	2014-06-24 2:45:00	164332	125504	1129161600000-446
click	2014-07-23 13:05:00	164332	167233	1225600000000-268
view	2018-12-22 15:07:00	164332	323498	1153267200000-265
click	2005-10-22 5:30:00	164332	522783	1333843200000-495
view	2014-12-12 20:46:00	164332	351883	1133636800000-147
click	2014-01-24 6:30:00	164332	438491	1187843200000-31
view	2018-02-28 2:43:00	164332	814873	1133296000000-171
click	2012-09-25 13:15:00	164332	544827	1206835200000-297
view	2018-12-22 2:13:00	164332	768588	1305244800000-164
click	2007-03-17 20:35:00	164332	139879	1151366400000-389



Commodity_Cox	Commodity_Des	Country_Code	Country_Name	Market_Year
1.	813300 Meal, Cottonseed	BR	Brazil	1969
2.	813300 Meal, Cottonseed	BR	Brazil	1969
3.	813300 Meal, Cottonseed	BR	Brazil	1969
4.	813300 Meal, Cottonseed	BR	Brazil	1969
5.	813300 Meal, Cottonseed	BR	Brazil	1969
6.	813300 Meal, Cottonseed	BR	Brazil	1969
7.	813300 Meal, Cottonseed	BR	Brazil	1969
8.	813300 Meal, Cottonseed	BR	Brazil	1969
9.	813300 Meal, Cottonseed	BR	Brazil	1969
10.	813300 Meal, Cottonseed	BR	Brazil	1969
11.	813300 Meal, Cottonseed	BR	Brazil	1969
12.	813300 Meal, Cottonseed	BR	Brazil	1969
13.	813300 Meal, Cottonseed	BR	Brazil	1969
14.	813300 Meal, Cottonseed	BR	Brazil	1969





 OPEN

 Refine

Data preparation: Cleaning and Wrangling

- Data Cleaning and Preparation
 - Handling Missing Data
 - Filtering Out Missing Data
 - Filling In Missing Data
 - Data Transformation
 - Removing Duplicates
 - Transforming Data Using a Function or Mapping
 - Replacing Values
 - Renaming Axis Indexes
 - Discretization and Binning
 - Detecting and Filtering Outliers
 - Permutation and Random Sampling
 - String Manipulation
 - String Object Methods
 - Regular Expressions



Data preparation: Cleaning and Wrangling

- Data Cleaning and Preparation

- Handling Missing Data

- Filtering Out Missing Data
- Filling In Missing Data

- Data Transformation

- Removing Duplicates
- Transforming Data Using a Function or Mapping
- Replacing Values
- Renaming Axis Indexes
- Discretization and Binning
- Detecting and Filtering Outliers
- Permutation and Random Sampling

- String Manipulation

- String Object Methods
- Regular Expressions



Data preparation: Cleaning and Wrangling

- Handling Missing Data
 - For numeric data, pandas uses the floating-point value **NaN (Not a Number)** to represent missing data

```
In [10]: string_data = pd.Series(['aardvark', 'artichoke', np.nan, 'avocado'])
```

```
In [11]: string_data
```

```
Out[11]:
```

```
0 aardvark
```

```
1 artichoke
```

```
2 NaN
```

```
3 avocado
```

```
dtype: object
```

```
In [12]: string_data.isnull()
```

```
Out[12]:
```

```
0 False
```

```
1 False
```

```
2 True
```

```
3 False
```

```
dtype: bool
```

Data preparation: Cleaning and Wrangling

- Handling Missing Data
 - The built-in Python **None** value is also treated as NA in object arrays

```
In [13]: string_data[0] = None
```

```
In [14]: string_data.isnull()
```

```
Out[14]:
```

```
0 True
```

```
1 False
```

```
2 True
```

```
3 False
```

```
dtype: bool
```

- NA Handling methods
 - **dropna** : Filter axis labels based on whether values for each label have missing data, with varying thresholds for how much missing data to tolerate
 - **fillna** : Fill in missing data with some value or using an interpolation method such as 'ffill' or 'bfill'
 - **isnull** : Return boolean values indicating which values are missing/NA
 - **notnull** : Negation of isnull

Data preparation: Cleaning and Wrangling

- Handling Missing Data: Filtering Out Missing Data
 - **dropna** on a Series, it returns the Series with only the non-null data and index values

```
In [15]: from numpy import nan as NA
```

```
In [16]: data = pd.Series([1, NA, 3.5, NA, 7])
```

```
In [17]: data.dropna()
```

```
Out[17]:
```

```
0 1.0
```

```
2 3.5
```

```
4 7.0
```

```
dtype: float64
```

This is equivalent to:

```
In [18]: data[data.notnull()]
```

```
Out[18]:
```

```
0 1.0
```

```
2 3.5
```

```
4 7.0
```

```
dtype: float64
```

Data preparation: Cleaning and Wrangling

- Handling Missing Data: Filtering Out Missing Data

- dropna** on Dataframe, it drops rows or columns that are all NA or only those containing any NAs. dropna by default drops any row containing a missing value:

- Passing how='all' will only drop rows that are all NA:

```
In [23]: data.dropna(how='all')
Out[23]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
3	NaN	6.5	3.0

- To drop columns, pass axis=1:

```
In [24]: data.dropna(axis=1, how='all')
```

```
In [19]: data = pd.DataFrame([[1., 6.5, 3.], [1., NA, NA],
.....:                        [NA, NA, NA], [NA, 6.5, 3.]])
```

```
In [20]: cleaned = data.dropna()
```

```
In [21]: data
```

```
Out[21]:
```

	0	1	2
0	1.0	6.5	3.0
1	1.0	NaN	NaN
2	NaN	NaN	NaN
3	NaN	6.5	3.0

```
In [22]: cleaned
```

```
Out[22]:
```

	0	1	2
0	1.0	6.5	3.0

Data preparation: Cleaning and Wrangling

- Handling Missing Data: Filtering Out Missing Data
 - Filter out DataFrame rows tends to concern time series data
 - Suppose you want to keep only rows containing a certain number of observations. You can indicate this with the `thresh` argument:

```
In [30]: df
Out[30]:
```

	0	1	2
0	-0.204708	NaN	NaN
1	-0.555730	NaN	NaN
2	0.092908	NaN	0.769023
3	1.246435	NaN	-1.296221
4	0.274992	0.228913	1.352917
5	0.886429	-2.001637	-0.371843
6	1.669025	-0.438570	-0.539741

```
In [31]: df.dropna()
Out[31]:
```

	0	1	2
4	0.274992	0.228913	1.352917
5	0.886429	-2.001637	-0.371843
6	1.669025	-0.438570	-0.539741

```
In [32]: df.dropna(thresh=2)
Out[32]:
```

	0	1	2
2	0.092908	NaN	0.769023
3	1.246435	NaN	-1.296221
4	0.274992	0.228913	1.352917
5	0.886429	-2.001637	-0.371843
6	1.669025	-0.438570	-0.539741

Data preparation: Cleaning and Wrangling

- Handling Missing Data: Filling In Missing Data
 - Rather than filtering out missing data (and potentially discarding other data along with it), you may want to fill in the “holes” in any number of ways
 - **fillna** method with a constant replaces missing values with that value
 - **fillna** with a dict, can use a different fill value for each column
 - **fillna** returns a new object, but it is possible to modify the existing object in-place

```
In [35]: _ = df.fillna(0, inplace=True)
```

```
In [33]: df.fillna(0)
Out[33]:
```

	0	1	2
0	-0.204708	0.000000	0.000000
1	-0.555730	0.000000	0.000000
2	0.092908	0.000000	0.769023
3	1.246435	0.000000	-1.296221
4	0.274992	0.228913	1.352917
5	0.886429	-2.001637	-0.371843
6	1.669025	-0.438570	-0.539741

```
In [34]: df.fillna({1: 0.5, 2: 0})
Out[34]:
```

	0	1	2
0	-0.204708	0.500000	0.000000
1	-0.555730	0.500000	0.000000
2	0.092908	0.500000	0.769023
3	1.246435	0.500000	-1.296221
4	0.274992	0.228913	1.352917
5	0.886429	-2.001637	-0.371843
6	1.669025	-0.438570	-0.539741

Data preparation: Cleaning and Wrangling

• Handling Missing Data: Filling In Missing Data

- A set of interpolation methods can be also used with **fillna**:

- **ffill**: to fill from front values

```
In [40]: df
Out[40]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	NaN	1.343810
3	-0.713544	NaN	-2.370232
4	-1.860761	NaN	NaN
5	-1.265934	NaN	NaN

- **bfill**: to backward fill the missing values

```
In [41]: df.fillna(method='ffill')
Out[41]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	0.124121	1.343810
3	-0.713544	0.124121	-2.370232
4	-1.860761	0.124121	-2.370232
5	-1.265934	0.124121	-2.370232

```
In [42]: df.fillna(method='ffill', limit=2)
Out[42]:
```

	0	1	2
0	0.476985	3.248944	-1.021228
1	-0.577087	0.124121	0.302614
2	0.523772	0.124121	1.343810
3	-0.713544	0.124121	-2.370232
4	-1.860761	NaN	-2.370232
5	-1.265934	NaN	-2.370232

- pass the mean or median value of a Series

Argument Description

value	Scalar value or dict-like object to use to fill missing values
method	Interpolation; by default 'ffill' if function called with no other arguments
axis	Axis to fill on; default axis=0
inplace	Modify the calling object without producing a copy
limit	For forward and backward filling, maximum number of consecutive periods to fill

```
In [43]: data = pd.Series([1., NA, 3.5, NA, 7])
```

```
In [44]: data.fillna(data.mean())
Out[44]:
```

0	1.000000
1	3.833333
2	3.500000
3	3.833333
4	7.000000

dtype: float64

Data preparation: Cleaning and Wrangling

- Data Cleaning and Preparation

- Handling Missing Data

- Filtering Out Missing Data
- Filling In Missing Data

- Data Transformation

- Removing Duplicates
- Transforming Data Using a Function or Mapping
- Replacing Values
- Renaming Axis Indexes
- Discretization and Binning
- Detecting and Filtering Outliers
- Permutation and Random Sampling

- String Manipulation

- String Object Methods
- Regular Expressions




Data preparation: Cleaning and Wrangling

- Data Transformation: Removing Duplicates
 - DataFrame method **duplicated** returns a boolean Series indicating whether each row is a duplicate
 - **drop_duplicates** solves this issue
 - Specify any subset of columns to detect duplicates

```
In [46]: data
Out[46]:
```

	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4
6	two	4




```
In [47]: data.duplicated()
Out[47]:
```

0	False
1	False
2	False
3	False
4	False
5	False
6	True

dtype: bool

```
In [48]: data.drop_duplicates()
Out[48]:
```


	k1	k2
0	one	1
1	two	1
2	one	2
3	two	3
4	one	3
5	two	4



```
In [49]: data['v1'] = range(7)
```

```
In [51]: data.drop_duplicates(['k1', 'k2'], keep='last')
Out[51]:
```

	k1	k2	v1
0	one	1	0
1	two	1	1
2	one	2	2
3	two	3	3
4	one	3	4
6	two	4	6



Keep='last' will return the last
duplicate
Instead of the first one

Data preparation: Cleaning and Wrangling

- Data Transformation: Transforming Data Using a Function or Mapping
 - For many datasets, you may wish to perform some transformation based on the values in an array, Series, or column in a DataFrame
 - Consider the following data collected about various kinds of meat

```
In [52]: data = pd.DataFrame({'food': ['bacon', 'pulled pork', 'bacon',
....:                                'Pastrami', 'corned beef', 'Bacon',
....:                                'pastrami', 'honey ham', 'nova lox'],
....:                        'ounces': [4, 3, 12, 6, 7.5, 8, 3, 5, 6]})
```

```
In [53]: data
Out[53]:
```

	food	ounces
0	bacon	4.0
1	pulled pork	3.0
2	bacon	12.0
3	Pastrami	6.0
4	corned beef	7.5
5	Bacon	8.0
6	pastrami	3.0
7	honey ham	5.0
8	nova lox	6.0

We want to add a column indicating the type of animal that each food came from. Let's write down a mapping of each distinct meat type to the kind of animal:

```
meat_to_animal = {
    'bacon': 'pig',
    'pulled pork': 'pig',
    'pastrami': 'cow',
    'corned beef': 'cow',
    'honey ham': 'pig',
    'nova lox': 'salmon'
}
```

```
In [55]: lowercased = data['food'].str.lower()
```

```
In [56]: lowercased
Out[56]:
```

	lowercased
0	bacon
1	pulled pork
2	bacon
3	pastrami
4	corned beef
5	bacon
6	pastrami
7	honey ham
8	nova lox

Name: food, dtype: object

Then we map

```
In [57]: data['animal'] = lowercased.map(meat_to_animal)
```

```
In [58]: data
Out[58]:
```

	food	ounces	animal
0	bacon	4.0	pig
1	pulled pork	3.0	pig
2	bacon	12.0	pig
3	Pastrami	6.0	cow
4	corned beef	7.5	cow
5	Bacon	8.0	pig
6	pastrami	3.0	cow
7	honey ham	5.0	pig
8	nova lox	6.0	salmon

Alternatively, we may use a lambda function to do the same

```
In [59]: data['food'].map(lambda x: meat_to_animal[x.lower()])
```

Data preparation: Cleaning and Wrangling

- Data Transformation: Replacing Values

- **replace** provides a simpler and more flexible way to substitute values

Let's consider this Series

- We can use `replace`, producing a new Series (unless you pass `inplace=True`)

```
In [62]: data.replace(-999, np.nan)
```

- **replace** multiple values at once, you instead pass a list and then the substitute value

```
In [63]: data.replace([-999, -1000], np.nan)
```

- To use a different replacement for each value, pass a list of substitutes

```
In [64]: data.replace([-999, -1000], [np.nan, 0])
```

- The argument passed can also be a dict

```
In [65]: data.replace({-999: np.nan, -1000: 0})
```

```
In [61]: data
Out[61]:
0      1.0
1    -999.0
2      2.0
3    -999.0
4   -1000.0
5      3.0
dtype: float64
```

```
Out[65]:
0      1.0
1     NaN
2      2.0
3     NaN
4      0.0
5      3.0
dtype: float64
```

Data preparation: Cleaning and Wrangling

- Data Transformation: Renaming Axis Indexes
 - Axis labels can be similarly transformed by a function or mapping of some form to produce new differently labelled objects

```
In [66]: data = pd.DataFrame(np.arange(12).reshape((3, 4)),  
.....:                      index=['Ohio', 'Colorado', 'New York'],  
.....:                      columns=['one', 'two', 'three', 'four'])
```

- Axis indexes have a map method

```
In [67]: transform = lambda x: x[:4].upper()
```

```
In [69]: data.index = data.index.map(transform)
```

```
In [70]: data
```

```
Out[70]:
```

	one	two	three	four
OHIO	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

Data preparation: Cleaning and Wrangling

- Data Transformation: Renaming Axis Indexes
 - **rename** method: to create a transformed version of a dataset without modifying the original

```
In [71]: data.rename(index=str.title, columns=str.upper)
Out[71]:
```

	ONE	TWO	THREE	FOUR
Ohio	0	1	2	3
Colo	4	5	6	7
New	8	9	10	11

- **rename** can be used in conjunction with a dict-like object providing new values for a subset of the axis labels

To modify a dataset in-place,
pass **inplace=True**

```
In [73]: data.rename(index={'OHIO': 'INDIANA'}, inplace=True)
```

```
In [72]: data.rename(index={'OHIO': 'INDIANA'},
....:                  columns={'three': 'peekaboo'})
Out[72]:
```

	one	two	peekaboo	four
INDIANA	0	1	2	3
COLO	4	5	6	7
NEW	8	9	10	11

Data preparation: Cleaning and Wrangling

- Data Transformation: Discretization and Binning
 - Continuous data is often discretized or otherwise separated into “bins” for analysis
- Let's divide these into bins of 18 to 25, 26 to 35, 36 to 60, and finally 61 and older: use **cut** function in pandas

```
In [75]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

```
In [76]: bins = [18, 25, 35, 60, 100]

In [77]: cats = pd.cut(ages, bins)

In [78]: cats
Out[78]:
[(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (35,
 60], (35, 60], (25, 35]]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

special
Categorical
object

Notación	Intervalo
$[a,b]$	$a \leq x \leq b$
$[a,b)$	$a \leq x < b$
$(a,b]$	$a < x \leq b$
(a,b)	$a < x < b$

Data preparation: Cleaning and Wrangling

Notación	Intervalo
$[a, b]$	$a \leq x \leq b$
$[a, b)$	$a \leq x < b$
$(a, b]$	$a < x \leq b$
(a, b)	$a < x < b$

- Data Transformation: Discretization and Binning
 - **pandas.cut**:
 - like an array of strings indicating the bin name
 - internally it contains a **categories** array specifying the distinct category names along with a labeling for the ages data in the **codes** attribute

```

In [79]: cats.codes
Out[79]: array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)

In [80]: cats.categories
Out[80]:
IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]]
              closed='right',
              dtype='interval[int64]')

In [81]: pd.value_counts(cats)
Out[81]:
(18, 25]      5
(35, 60]      3
(25, 35]      3
(60, 100]     1
dtype: int64

```

Histogramming

a parenthesis means that the side is *open*, while the square bracket means it is *closed* (inclusive);

It can be changed by passing `right=False`

```
In [82]: pd.cut(ages, [18, 26, 36, 61, 100], right=False)
```

Data preparation: Cleaning and Wrangling

- Data Transformation:
Discretization and Binning

- **pandas.cut:**

- It allows to pass own bin names by passing a list or array to the labels option

```
In [83]: group_names = ['Youth', 'YoungAdult', 'MiddleAged', 'Senior']

In [84]: pd.cut(ages, bins, labels=group_names)
Out[84]:
[Youth, Youth, Youth, YoungAdult, Youth, ..., YoungAdult, Senior, MiddleAged, MiddleAged, YoungAdult]
Length: 12
Categories (4, object): [Youth < YoungAdult < MiddleAged < Senior]
```

- Compute equal-length bins based on the minimum and maximum values in the data

precision=2 option limits the decimal precision to two digits

```
In [85]: data = np.random.rand(20)

In [86]: pd.cut(data, 4, precision=2)
Out[86]:
[(0.34, 0.55], (0.34, 0.55], (0.76, 0.97], (0.76, 0.97], (0.34, 0.55], ..., (0.34, 0.55], (0.34, 0.55], (0.55, 0.76], (0.34, 0.55], (0.12, 0.34]]
Length: 20
Categories (4, interval[float64]): [(0.12, 0.34] < (0.34, 0.55] < (0.55, 0.76] < (0.76, 0.97]]
```

Data preparation: Cleaning and Wrangling


- Data Transformation: Discretization and Binning
 - **pandas.qcut**: bins the data based on sample **quantiles**
 - by definition roughly equal-size bins are obtained

```

In [87]: data = np.random.randn(1000) # Normally distributed

In [88]: cats = pd.qcut(data, 4) # Cut into quartiles

In [89]: cats
Out[89]:
[(-0.0265, 0.62], (0.62, 3.928], (-0.68, -0.0265], (0.62, 3.928], (-0.0265, 0.62]
, ..., (-0.68, -0.0265], (-0.68, -0.0265], (-2.95, -0.68], (0.62, 3.928], (-0.68,
-0.0265]]
Length: 1000
Categories (4, interval[float64]): [(-2.95, -0.68] < (-0.68, -0.0265] < (-0.0265,
0.62] <
(0.62, 3.928]]
  
```



```

In [90]: pd.value_counts(cats)
Out[90]:
(0.62, 3.928]      250
(-0.0265, 0.62]    250
(-0.68, -0.0265]    250
(-2.95, -0.68]      250
dtype: int64
  
```

- Similar to cut you can pass your own quantiles (numbers between 0 and 1, inclusive)

```
In [91]: pd.qcut(data, [0, 0.1, 0.5, 0.9, 1.])
```

Data preparation: Cleaning and Wrangling

• Data Transformation: Detecting and Filtering Outliers

- It's a matter of array indexing
- Having this DataFrame:
 - Let's find values in one of the columns exceeding 3 in absolute value:

```

In [94]: col = data[2]

In [95]: col[np.abs(col) > 3]
Out[95]:
41    -3.399312
136    -3.745356
Name: 2, dtype: float64

```

- cap values outside the interval -3 to 3:

np.sign(data) produces 1 and -1
 values based on whether the values
 in data are positive or negative

```

In [92]: data = pd.DataFrame(np.random.randn(1000, 4))

In [93]: data.describe()
Out[93]:

```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.049091	0.026112	-0.002544	-0.051827
std	0.996947	1.007458	0.995232	0.998311
min	-3.645860	-3.184377	-3.745356	-3.428254
25%	-0.599807	-0.612162	-0.687373	-0.747478
50%	0.047101	-0.013609	-0.022158	-0.088274
75%	0.756646	0.695298	0.699046	0.623331
max	2.653656	3.525865	2.735527	3.366626

```

In [97]: data[np.abs(data) > 3] = np.sign(data) * 3

In [98]: data.describe()
Out[98]:

```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	0.050286	0.025567	-0.001399	-0.051765
std	0.992920	1.004214	0.991414	0.995761
min	-3.000000	-3.000000	-3.000000	-3.000000
25%	-0.599807	-0.612162	-0.687373	-0.747478
50%	0.047101	-0.013609	-0.022158	-0.088274
75%	0.756646	0.695298	0.699046	0.623331
max	2.653656	3.000000	2.735527	3.000000

Data preparation: Cleaning and Wrangling

• Data Transformation: Detecting and Filtering Outliers

- Remove all the random numbers that lie in the lowest **quantile** and the highest quantile

```
size=200
```

```
x = pd.Series(np.random.normal(size=size)) # 200 values
x = x[x.between(x.quantile(.15), x.quantile(.85))] # without outliers
```

```
print(x) # Now only 140 values
```

```

0      0.691716
1      0.519569
3     -0.145321
4     -0.490216
5     -0.357589
...
191    -0.254481
192     0.808355
194    -0.755343
196    -0.132525
199    -0.494869
Length: 140, dtype: float64

```

- With two axes

```

import datetime

todays_date = datetime.datetime.now().date()
dates = pd.date_range(todays_date-datetime.timedelta(10), periods=size, freq='D')

rando_nums = np.random.normal(size=size)
columns = ['rando']

df = pd.DataFrame(rando_nums, index=dates, columns=columns)
df      # 200 random numbers indexed by days in a week
df.plot().get_figure()

```


```

y = df['rando']
removed_outliers = y.between(y.quantile(.05), y.quantile(.95))

print(str(y[removed_outliers].size) + "/" + str(size) + " data points remain.")

y[removed_outliers].plot().get_figure()

```



- np.percentile**

Data preparation: Cleaning and Wrangling

- Data Transformation: Permutation and Random Sampling
 - **permutation** with the length of the axis you want to permute produces an array of integers indicating the new ordering

```
In [100]: df = pd.DataFrame(np.arange(5 * 4).reshape((5, 4)))
```

```
In [101]: sampler = np.random.permutation(5)
```

```
In [102]: sampler
```

```
Out[102]: array([3, 1, 4, 2, 0])
```

```
In [103]: df
```

```
Out[103]:
```

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

- To select a random subset without replacement, you can use the **sample** method on Series and DataFrame

```
In [105]: df.sample(n=3)
```

```
Out[105]:
```

	0	1	2	3
3	12	13	14	15
4	16	17	18	19
2	8	9	10	11

```
In [104]: df.take(sampler)
```

```
Out[104]:
```

	0	1	2	3
3	12	13	14	15
1	4	5	6	7
4	16	17	18	19
2	8	9	10	11
0	0	1	2	3

Data preparation: Cleaning and Wrangling

- Data Cleaning and Preparation
 - Handling Missing Data
 - Filtering Out Missing Data
 - Filling In Missing Data
 - Data Transformation
 - Removing Duplicates
 - Transforming Data Using a Function or Mapping
 - Replacing Values
 - Renaming Axis Indexes
 - Discretization and Binning
 - Detecting and Filtering Outliers
 - Permutation and Random Sampling
 - String Manipulation
 - String Object Methods
 - Regular Expressions



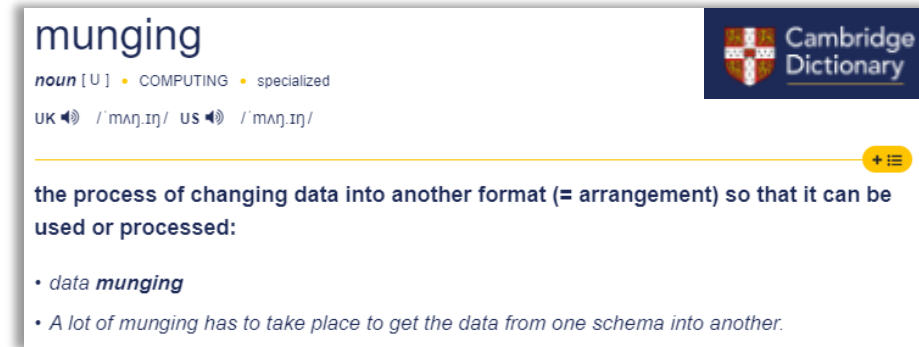
Data preparation: Cleaning and Wrangling

- String Manipulation: String Object Methods
 - In many **string munging** and scripting applications, built-in string methods are sufficient.
Example

- a comma-separated string can be broken into pieces with split

```
In [134]: val = 'a,b, guido'

In [135]: val.split(',')
Out[135]: ['a', 'b', ' guido']
```



- split is often combined with strip to trim whitespace (including line breaks)

```
In [136]: pieces = [x.strip() for x in val.split(',')]

In [137]: pieces
Out[137]: ['a', 'b', 'guido']
```

- These substrings could be concatenated together with a two-colon delimiter by passing a list or tuple to the **join** method on the string '::'

```
In [140]: '::'.join(pieces)
Out[140]: 'a::b::guido'
```

Data preparation: Cleaning and Wrangling

- String Manipulation: String Object Methods
 - Other methods are concerned with locating substrings
 - **index** and **find** methods to detect substrings
 - **count** returns the number of occurrences of a particular substring
 - **replace** will substitute occurrences of one pattern for another.

```
In [145]: val.count(',')  
Out[145]: 2
```

```
In [146]: val.replace(',', '::  
Out[146]: 'a::b:: guido'  
  
In [147]: val.replace(',', '')  
Out[147]: 'ab guido'
```

It is commonly used to delete patterns, too, by passing an empty string

```
In [134]: val = 'a,b, guido'
```

```
In [141]: 'guido' in val  
Out[141]: True
```

```
In [142]: val.index(',')  
Out[142]: 1
```

```
In [143]: val.find(':')  
Out[143]: -1
```

index raises an exception if the string isn't found (versus returning -1)

`val.index(':')`

Data preparation: Cleaning and Wrangling

- String Manipulation: String Object Methods
 - Python built-in string methods

Argument	Description
<code>count</code>	Return the number of non-overlapping occurrences of substring in the string.
<code>endswith</code>	Returns <code>True</code> if string ends with suffix.
<code>startswith</code>	Returns <code>True</code> if string starts with prefix.
<code>join</code>	Use string as delimiter for concatenating a sequence of other strings.
<code>index</code>	Return position of first character in substring if found in the string; raises <code>ValueError</code> if not found.
<code>find</code>	Return position of first character of <i>first</i> occurrence of substring in the string; like <code>index</code> , but returns <code>-1</code> if not found.
<code>rfind</code>	Return position of first character of <i>last</i> occurrence of substring in the string; returns <code>-1</code> if not found.
<code>replace</code>	Replace occurrences of string with another string.
<code>strip</code> , <code>rstrip</code> , <code>lstrip</code>	Trim whitespace, including newlines; equivalent to <code>x.strip()</code> (and <code>rstrip</code> , <code>lstrip</code> , respectively) for each element.
<code>split</code>	Break string into list of substrings using passed delimiter.
<code>lower</code>	Convert alphabet characters to lowercase.
<code>upper</code>	Convert alphabet characters to uppercase.
<code>casefold</code>	Convert characters to lowercase, and convert any region-specific variable character combinations to a common comparable form.
<code>ljust</code> , <code>rjust</code>	Left justify or right justify, respectively; pad opposite side of string with spaces (or some other fill character) to return a string with a minimum width.

Data preparation: Cleaning and Wrangling

- String Manipulation: Regular Expressions

- A single expression, commonly called a **regex**, is a string formed according to the regular expression language

```
In [148]: import re
```

- Python's built-in **re** module is responsible for applying regular expressions to strings

- **re** functions fall into three categories

- Pattern matching
- Substitution
- Splitting

Describing one or more
whitespace characters
is `\s+`

```
In [149]: text = "foo    bar\t baz  \tqux"

In [150]: re.split('\s+', text)
Out[150]: ['foo', 'bar', 'baz', 'qux']
```

```
In [151]: regex = re.compile('\s+')

In [152]: regex.split(text)
Out[152]: ['foo', 'bar', 'baz', 'qux']
```

Compiling by
yourself

To avoid unwanted escaping with `\` in a regular expression, use *raw* string literals like `r'C:\x'` instead of the equivalent `'C:\\x'`.

```
In [153]: regex.findall(text)
Out[153]: [' ', '\t', ' ', '\t']
```

Data preparation: Cleaning and Wrangling

- String Manipulation: Regular Expressions
 - **findall** returns all matches in a string
 - **search** returns only the first match
 - **match** only matches at the beginning of the string

```
text = """Dave dave@google.com
Steve steve@gmail.com
Rob rob@gmail.com
Ryan ryan@yahoo.com
"""
pattern = r'[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,4}'

# re.IGNORECASE makes the regex case-insensitive
regex = re.compile(pattern, flags=re.IGNORECASE)
```

```
In [155]: regex.findall(text)
Out[155]:
['dave@google.com',
 'steve@gmail.com',
 'rob@gmail.com',
 'ryan@yahoo.com']
```

```
In [156]: m = regex.search(text)
```

```
In [157]: m
```

```
Out[157]: <_sre.SRE_Match object; span=(5, 20), match='dave@google.com'>
```

```
In [158]: text[m.start():m.end()]
```

```
Out[158]: 'dave@google.com'
```

None, because it only will match if the
pattern occurs at the start of the string

```
In [159]: print(regex.match(text))
None
```


Data preparation: Cleaning and Wrangling

- String Manipulation: Regular Expressions

- sub** will return a new string with occurrences of the pattern replaced by the a new string

```
In [160]: print(regex.sub('REDACTED', text))
Dave REDACTED
Steve REDACTED
Rob REDACTED
Ryan REDACTED
```

- To segment components (username, domain name, and domain suffix) put parentheses around the parts of the pattern

```
In [161]: pattern = r'([A-Z0-9._%+-]+)@([A-Z0-9.-]+)\.([A-Z]{2,4})'

In [162]: regex = re.compile(pattern, flags=re.IGNORECASE)
```

- A match object produced by this modified regex returns a tuple of the pattern components with its **groups** method

```
text = """Dave dave@google.com
Steve steve@gmail.com
Rob rob@gmail.com
Ryan ryan@yahoo.com
"""

pattern = r'[A-Z0-9._%+-]+@[A-Z0-9.-]+\.[A-Z]{2,4}'

# re.IGNORECASE makes the regex case-insensitive
regex = re.compile(pattern, flags=re.IGNORECASE)
```

```
In [163]: m = regex.match('wesm@bright.net')

In [164]: m.groups()
Out[164]: ('wesm', 'bright', 'net')
```

Data preparation: Cleaning and Wrangling

- String Manipulation: Regular Expressions

- findall** returns a list of tuples when the pattern has groups

```
In [165]: regex.findall(text)
Out[165]:
[('dave', 'google', 'com'),
 ('steve', 'gmail', 'com'),
 ('rob', 'gmail', 'com'),
 ('ryan', 'yahoo', 'com')]
```

- sub** also has access to groups in each match using special symbols like \1 and \2
 - The symbol \1 corresponds to the first matched group, \2 corresponds to the second, and so forth

```
In [166]: print(regex.sub(r'Username: \1, Domain: \2, Suffix: \3', text))
Dave Username: dave, Domain: google, Suffix: com
Steve Username: steve, Domain: gmail, Suffix: com
Rob Username: rob, Domain: gmail, Suffix: com
Ryan Username: ryan, Domain: yahoo, Suffix: com
```